
Patterns for Parallel Programming

Timothy G. Mattson

Beverly A. Sanders

Berna L. Massingill

◆◆Addison-Wesley

Boston • San Francisco • New York • Toronto • Montreal
London • Munich • Paris • Madrid
Capetown • Sydney • Tokyo • Singapore • Mexico City

Contents

Preface	x
1 A Pattern Language for Parallel Programming	1
1.1 Introduction	1
1.2 Parallel Programming	3
1.3 Design Patterns and Pattern Languages	4
1.4 A Pattern Language for Parallel Programming	5
2 Background and Jargon of Parallel Computing	7
2.1 Concurrency in Parallel Programs Versus Operating Systems	7
2.2 Parallel Architectures: A Brief Introduction	8
2.2.1 Flynn's Taxonomy	8
2.2.2 A Further Breakdown of MIMD	9
2.2.3 Summary	12
2.3 Parallel Programming Environments	12
2.4 The Jargon of Parallel Computing	16
2.5 A Quantitative Look at Parallel Computation	18
2.6 Communication	21
2.6.1 Latency and Bandwidth	21
2.6.2 Overlapping Communication and Computation and Latency Hiding	22
2.7 Summary	23
3 The <i>Finding Concurrency</i> Design Space	24
3.1 About the Design Space	24
3.1.1 Overview	25
3.1.2 Using the Decomposition Patterns	26
3.1.3 Background for Examples	26
3.2 The <i>Task Decomposition</i> Pattern	29
3.3 The <i>Data Decomposition</i> Pattern	34
3.4 The <i>Group Tasks</i> Pattern	39
3.5 The <i>Order Tasks</i> Pattern	42
3.6 The <i>Data Sharing</i> Pattern	44
3.7 The <i>Design Evaluation</i> Pattern	49
3.8 Summary	55
4 The <i>Algorithm Structure</i> Design Space	57
4.1 Introduction	57
4.2 Choosing an <i>Algorithm Structure</i> Pattern	59
4.2.1 Target Platform	59

4.2.2	Major Organizing Principle	60
4.2.3	The <i>Algorithm Structure</i> Decision Tree	60
4.2.4	Re-evaluation	62
4.3	Examples	62
4.3.1	Medical Imaging	62
4.3.2	Molecular Dynamics	63
4.4	The <i>Task Parallelism</i> Pattern	64
4.5	The <i>Divide and Conquer</i> Pattern	73
4.6	The <i>Geometric Decomposition</i> Pattern	79
4.7	The <i>Recursive Data</i> Pattern	97
4.8	The <i>Pipeline</i> Pattern	103
4.9	The <i>Event-Based Coordination</i> Pattern	114
5	The <i>Supporting Structures</i> Design Space	121
5.1	Introduction	121
5.1.1	Program Structuring Patterns	122
5.1.2	Patterns Representing Data Structures	123
5.2	Forces	123
5.3	Choosing the Patterns	125
5.4	The <i>SPMD</i> Pattern	126
5.5	The <i>Master/Worker</i> Pattern	143
5.6	The <i>Loop Parallelism</i> Pattern	152
5.7	The <i>Fork/Join</i> Pattern	167
5.8	The <i>Shared Data</i> Pattern	173
5.9	The <i>Shared Queue</i> Pattern	183
5.10	The <i>Distributed Array</i> Pattern	198
5.11	Other Supporting Structures	211
5.11.1	SIMD	211
5.11.2	MPMD	212
5.11.3	Client-Server Computing	214
5.11.4	Concurrent Programming with Declarative Languages	214
5.11.5	Problem-Solving Environments	215
6	The <i>Implementation Mechanisms</i> Design Space	216
6.1	Overview	217
6.2	UE Management	217
6.2.1	Thread Creation/Destruction	218
6.2.2	Process Creation/Destruction	220
6.3	Synchronization	221
6.3.1	Memory Synchronization and Fences	221
6.3.2	Barriers	226
6.3.3	Mutual Exclusion	229
6.4	Communication	237
6.4.1	Message Passing	238
6.4.2	Collective Communication	245
6.4.3	Other Communication Constructs	251

Appendix A A Brief Introduction to OpenMP	253
A.1 Core Concepts	254
A.2 Structured Blocks and Directive Formats	257
A.3 Worksharing	259
A.4 Data Environment Clauses	262
A.5 The OpenMP Runtime Library	265
A.6 Synchronization	266
A.7 The Schedule Clause	270
A.8 The Rest of the Language	272
Appendix B A Brief Introduction to MPI	273
B.1 Concepts	273
B.2 Getting Started	275
B.3 Basic Point-to-Point Message Passing	277
B.4 Collective Operations	279
B.5 Advanced Point-to-Point Message Passing	283
B.6 MPI and Fortran	288
B.7 Conclusion	290
Appendix C A Brief Introduction to Concurrent Programming in Java	291
C.1 Creating Threads	293
C.2 Atomicity, Memory Synchronization, and the volatile Keyword . .	297
C.3 Synchronized Blocks	297
C.4 Wait and Notify	299
C.5 Locks	301
C.6 Other Synchronization Mechanisms and Shared Data Structures . .	303
C.7 Interrupts	304
Glossary	307
Bibliography	317
About the Authors	333
Index	335